

# Regularization

May 28, 2012

## 1 The problem of overfitting.

- **Underfitting** - High bias - doesn't fit the data very well
- **Overfitting** - High variance - high order polynomial, fits data very well, space of possible hypothesis is too variable to give a good one
  - If we have too many features, the learned hypothesis may fit the training set very well ( $J(\Theta) \approx 0$ ), but fail to generalize to new examples (predict prices on new examples)

### 1.1 Addressing overfitting

Lots of features and little training data, overfitting can occur.  
If we think overfitting is occurring, what can we do?

- Plotting hypothesis - to select an appropriate degree polynomial. Doesn't always work, you can have too many features.
- Reduce number of features.
  - Manually select which features to keep.
  - Model selection algorithm
- Regularization.
  - Keep all features, but reduce magnitude/values of parameters  $\Theta_j$ .
  - Works well when we have a lot of features, each of which contributes a bit to predicting  $y$ .

## 2 Cost function

Take parameters:

$\Theta_0 + \Theta_1 x + \Theta_2 x^2 \rightarrow$  just right

$\Theta_0 + \Theta_1 x + \Theta_2 x^2 + \Theta_3 x^3 + \Theta_4 x^4 \rightarrow$  overfitted

Suppose we penalize and make  $\Theta_3, \Theta_4$  really small.

$$\min_{\Theta} \frac{1}{2m} \sum_{i=1}^m (h_{\Theta}(x^{(i)}) - y^{(i)})^2 + 1000\Theta_3^2 + 1000\Theta_4^2$$

The only way to make this cost function small is when  $\Theta_3 \approx 0$  and  $\Theta_4 \approx 0$  and thus essentially getting rid of the high order terms.

In general:

Small values for parameters  $\Theta_0, \Theta_1, \dots, \Theta_n$

- “Simpler” hypothesis
- Less prone to overfitting

Ex: Housing:

Features:  $x_1, x_2, \dots, x_{100}$  - hard to pick in advance which are the ones that are less likely to be relevant

Parameters:  $\Theta_0, \Theta_1, \dots, \Theta_{100}$

We’ll take the cost function and add a regularization term at the end to shrink all of the parameters. Not the  $\Theta_0$  term

$$J(\Theta) = \frac{1}{2m} \left[ \sum_{i=1}^m (h_{\Theta}(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n \Theta_j^2 \right]$$

$\lambda$ - regularization parameter, controls a tradeoff between two different goals:

- we would like to train the training set well (first term)
- we want to keep the parameters small (regularization term) and therefore keeping the hypothesis simple to avoid overfitting

What if  $\lambda$  is set to an extremely large value (perhaps too large for our problem, say  $\lambda = 10^{10}$ )?

We would end up penalizing all the parameters so they end up close to zero, ending up with a hypothesis:

$$h_{\Theta}(x) = \Theta_0$$

and underfit the training set. Hypothesis has too high bias or preconception.

## 3 Regularized linear regression

For linear regression we have two algorithms

- Gradient descent

- Normal equation

We'll generalize them for linear regression.

$$J(\Theta) = \frac{1}{2m} \left[ \sum_{i=1}^m (h_{\Theta}(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n \Theta_j^2 \right]$$

### 3.1 Gradient descent

$$\begin{aligned} & \text{Repeat} \{ \\ & \quad \Theta_j := \Theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\Theta}(x)^{(i)} - y^{(i)}) x_j^{(i)} \\ & \quad \quad (j = 0, 1, 2, 3, \dots, n) \\ & \} \end{aligned}$$

With regularization for linear regression:

$$\begin{aligned} & \text{Repeat} \{ \\ & \quad \Theta_0 := \Theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\Theta}(x)^{(i)} - y^{(i)}) x_0^{(i)} \\ & \quad \Theta_j := \Theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\Theta}(x)^{(i)} - y^{(i)}) x_j^{(i)} + \frac{\lambda}{m} \Theta_j \\ & \quad \quad (j = 1, 2, 3, \dots, n) \\ & \} \end{aligned}$$

The update for  $\Theta_j$  can be written equivalently as

$$\Theta_j := \Theta_j (1 - \alpha \frac{\lambda}{m}) - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\Theta}(x)^{(i)} - y^{(i)}) x_j^{(i)}$$

The first term,  $1 - \alpha \frac{\lambda}{m} < 1$ , has an effect of shrinking  $\Theta_j$  and then the second term is the same as original gradient descent update.

### 3.2 Normal equation

$$X = \begin{bmatrix} (x^{(1)})^\top \\ \vdots \\ (x^{(m)})^\top \end{bmatrix} \quad y = \begin{bmatrix} y^{(1)} \\ \vdots \\ y^{(m)} \end{bmatrix} \in \mathbb{R}^m$$

$$\Theta = (X^\top X + \lambda \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & \ddots & \vdots \\ 0 & 0 & 0 & \dots & 1 \end{bmatrix})^{-1} X^\top y$$

The inner matrix has dimensions:  $(n + 1) \times (n + 1)$

### 3.2.1 Non-invertibility

Suppose  $m \leq n$ , (# examples < #features)

pinv will give an example that kindof makes sense, numerically correct but not necessarily a good hypothesis.

if  $\lambda > 0$ , it is possible to prove that term of

$$\Theta = (X^T X + \lambda \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & \ddots & \vdots \\ 0 & 0 & 0 & \dots & 1 \end{bmatrix})^{-1} X^T y$$

in parentheses will be invertible.

## 4 Regularized Logistic Regression

$$J(\Theta) = - \left[ \frac{1}{m} \sum_{i=1}^m y^{(i)} \log h_{\Theta}(x^{(i)}) + (1 - y^{(i)}) \log(1 - h_{\Theta}(x^{(i)})) \right] + \lambda \sum_{j=1}^n \Theta_j^2$$

So now, even if we're fitting a very high order polynomial, we'll more likely get a decision boundary that looks reasonable.

Implementation:

*Repeat* {

$$\Theta_0 := \Theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\Theta}(x)^{(i)} - y^{(i)}) x_0^{(i)}$$

$$\Theta_j := \Theta_j - \alpha \left[ \frac{1}{m} \sum_{i=1}^m (h_{\Theta}(x)^{(i)} - y^{(i)}) x_j^{(i)} + \frac{\lambda}{m} \Theta_j \right]$$

}  $(j = 1, 2, 3, \dots, n)$

The difference from linear regression is that the hypothesis is different

$$h_{\Theta}(x) = \frac{1}{1 + e^{-\Theta^T x}}$$

### 4.1 Advanced optimization

function [jVal, gradient] = costFunction(theta)

% below code now must include the regularization parameter

jval = [code to compute  $J(\Theta)$ ];

```

gradient(1) = [code to compute  $\frac{\partial}{\partial \Theta_0} J(\Theta)$ ];
gradient(2) = [code to compute  $\frac{\partial}{\partial \Theta_1} J(\Theta)$ ];
⋮
gradient(n+1) = [code to compute  $\frac{\partial}{\partial \Theta_n} J(\Theta)$ ];

```

Then you pass the results to fminunc or other advanced optimization algorithm, the params you get out will correspond to regularized logistic regression.